

Faster Matrices

Aidan J. Nyquist
CERN-HSF, Google Summer of Code
(Dated: March 2018)

Proposal for "Faster matrix algebra for ATLAS" Google Summer of Code, CERN-HSF Project.
The purpose of this project is to increase the performance of ATLAS by optimizing its symmetric matrix operations. This will decrease computing and storage demands, and increase speed.

INTRODUCTION

The ATLAS Experiment uses Eigen extensively. However, Eigen handles symmetric matrices inefficiently. Since a large portion of ATLAS' algorithms are made up of symmetric matrices, this project is able to optimize a large niche in the code. By implementing a symmetric matrix class, the storage space, and the complexity of certain algorithms like addition of a dense square matrix ($n \times n$) originally growing at n^2 becomes $\frac{n(n+1)}{2}$. This approaches a 50% complexity and memory decrease for larger and larger matrices. Of course, not all algorithms become as efficient. For example symmetric matrices are not generally closed under matrix multiplication, so it would be unreasonable to always expect a dramatic speed increase. However, the memory decrease alone should be more than enough to justify a symmetric matrix class. The algorithm efficiency is just the frosting on the cake.

PROJECT GOALS

"A working implementation of symmetric matrices in Eigen, ready to be submitted as a patch for Eigen." - Proposal Page

High Priority

- Symmetric Matrix Class
- Basic matrix operations
- Seamless integration with Eigen Matrices ¹
- Use Eigen conventions and style in code
- Decrease in memory cost
- Tests
- Well documented

¹ There are *a lot* of operations Eigen can perform on matrices. By integrating with Eigen there will be no need to rewrite any algorithms that won't have a direct speedup from a Symmetric Matrix Class.

Medium Priority

- Increase in performance
 - Benchmarks to show improvement
- Eigenvalue algorithms
- Reusable code

Low Priority

- Various matrix decompositions
- Complex symmetric matrices
- Skew-symmetric matrices
- Sparse symmetric matrices

TIMELINE

The following is a rough outline. There is still some research and dialogue to be done before and during the summer. Many tasks might bleed back and forth through the timeline when new insights are found whilst the project progresses.

Pre-Summer

The best time to get a head start on the project is right now. One of the most important things I could do is to continue interacting with Eigen library and community. That way, I'll have practiced swimming before I have dive into Eigen.

I also will reboot my blog. This will give me a way of introducing myself and my project. I will also use it as communication platform for the rest of the summer and beyond. Meaning at minimum, there will be a blog post for every section outlined in this timeline.

There are many ways of extending Eigen, picking the right one first would be more than ideal. To make sure I pick the right one first, I will study current Eigen Matrix classes, prototype different implementations of classes and methods, and engage with the community and my mentors.

First Third of Summer

- Finish prototyping integration with Eigen
- Extend best prototype
- Implement most optimal data structure
- Implement as many functions as possible
 - Accessors
 - Arithmetic (+, +=, -, ·)
 - Comparison (==, !=)

Second Third of Summer

- Continue implementing functions ²
 - Inverse
 - REF / RREF
 - Eigenvalues & Eigenvectors
 - * Research
 - * QR decomposition
 - * Bidiagonalization
- Organize tests and write benchmarks
- Refactor any poor code
- If time allows, begin a low priority task or two

Last Third of Summer

- If time allows, finish a low priority task or two
- Finish any lingering tasks
- Prepare and submit patch
 - If the patch won't get merged, I will republish the code as a 3rd party plugin.

Post-Summer

If the patch does not initially get accepted, or the project runs into roadblocks, I plan to see the project through—whatever continuation of dialogue or programming that might involve.

Bugs and questions regarding my code are things I plan to watch for a few years. I also hope to be able to answer any questions that might arise on as many communication platforms as I can watch. Any bugs that reveal themselves later in my code, I also plan to fix.

² Some of these functions will need to be reviewed before being implemented, as there may be no advantage rewriting them.

RELATED WORK

ATLAS uses various symmetric matrices when computing charged particle tracks. It currently stores a lot of 5×5 covariant matrices, so an optimized symmetric matrix class would greatly decrease memory costs. It also uses the following matrix product for Kalman Filters: $A \cdot S \cdot A^T$, where S is symmetric and therefore the resulting matrix is also symmetric. Because the resulting matrix is known to be symmetric, the matrix multiplication can be optimized by only computing the upper triangle.

Current support for symmetric matrices in Eigen are in the form of views. While these views can support optimized algorithms, they do not optimize memory in any way. There are other popular libraries that have somewhat implemented symmetric matrices: the GNU Scientific Library, and Boost uBLAS. Studying the code of these libraries will be insightful for my project. so a part of the timeline might be taken therefor.

There is a lot of academic work related to symmetric matrices. While I plan to review all relevant information before the summer, some time might be spent during the summer re-analyzing these resources.

BIOGRAPHICAL INFORMATION

I am currently pursuing a B.S. in Mathematics and a B.A. in Physics and minoring in Computer Science. This project is a beautiful mixture of indirectly contributing to the physics community, by directly contributing to the programming community by using my mathematical abilities. I've taken a couple of couple of linear algebra courses, and have used my programming skills, for recreation, open-source contribution, and paid work for a few years now. I have prototyped some code for this project already: github.com/ajjony/symmat/ I feel very qualified to take up this project, not only because of my background and hard-skills, but as well as because of my relentless passion and drive for projects like these.

CONTACT INFORMATION

- Email: contact@aidannyquist.com
- University Email: ajnyquist@northpark.edu
- Website (WIP): aidannyquist.com
- IRC/GitHub: [ajjony](#)
- Phone: 971-237-5505
- Address (Legal/Summer): 3580 NW Hill Rd, McMinnville, OR 97128